

Let's say A and B play 7 games and A wins 4 of them. **What would be a reasonable estimate for A's win probability (over B)?** A natural estimate would be 4/7. More generally, if A wins  $x$  games, then the estimator  $\hat{p} = x/7$  is an unbiased estimator of the true win probability  $P$ . This is not difficult to show: since  $x \sim \text{Binom}(7, p)$ , we have

$$\mathbb{E}[x/7] = \mathbb{E}[x]/7 = (7p)/7 = p.$$

Now, let's say that A and B play a best-of-7 series (or equivalently, first to 4 wins), and A wins 4 of them. It seems natural to estimate A's win probability by 4/7 again. If we define the estimator

$$\hat{p} = \frac{\# \text{ games A wins}}{\# \text{ games}},$$

**does this estimator have any nice properties?** It's hard to say anything about this estimator (if someone knows something, please let me know!). In fact, it's not even unbiased!

### Estimator is biased: a simulation

First, let's set up two functions. The first function, `SampleSeries`, simulates a series which ends when either player 1 reaches `n_wins1` wins or when player 2 reaches `n_wins2` wins. (The true win probability for player 1 is  $p$ .) The second function, `EstimateWinProb`, samples `n_sim` series. For each series, it estimates player 1's win probability as in  $\hat{p}$  above, then returns the mean of  $\hat{p}$  over the `n_sim` simulations.


```
# Player 1 has win probability p. Simulates a series of games which
ends when
# player 1 reaches `n_wins1` wins or when player 2 reaches `n_wins2`
wins.
SampleSeries <- function(p, n_wins1, n_wins2 = n_wins1) {
  game_results <- rbinom(n_wins1 + n_wins2 - 1, size = 1, prob = p)
  if (sum(game_results) >= n_wins1) {
    n_games <- which(game_results == 1)[n_wins1]
  } else {
    n_games <- which(game_results == 0)[n_wins2]
  }
  player1_wins <- sum(game_results[1:n_games])
  return(c(player1_wins, n_games - player1_wins))
}

# Run SampleSeries `n_sim` times and for each run, estimate p by player
1's
# win percentage in that series. Returns the estimated probabilities.
EstimateWinProb <- function(p, n_sim, n_wins1, n_wins2 = n_wins1) {
  win_data <- replicate(n_sim, SampleSeries(p, n_wins1, n_wins2))
  prob_estimates <- apply(win_data, 2, function(x) x[1] / (x[1] +
x[2]))
  return(mean(prob_estimates))
}
```

Next, we run a Monte Carlo simulation (20,000 simulations) to compute the mean of  $\hat{p}$  for a range of values of  $P$ . (Because of the symmetry inherent in the setup, we only consider  $p \geq 0.5$

.)

```
set.seed(1)
p <- 20:40 / 40
n_wins1 <- 4
n_sim <- 20000
mean_phat <- sapply(p, function(p) EstimateWinProb(p, n_sim, n_wins1))
plot(p, mean_phat, asp = 1, pch = 16, ylab = "Win prob estimator mean",
     main = "First to 4 wins")
abline(0, 1)
```



If  $\hat{p}$  were unbiased, we would expect the black dots to lie on the black line  $y = x$ ; it's clear that they don't. In this setting,  $\hat{p}$  is biased upwards.

### **Estimator is biased: exact computation**

I was not able to find a nice closed-form expression for  $\mathbb{E}[\hat{p}]$ : for a best-of-7 series, [WolframAlpha](#) "simplified" my expression to give

$$\mathbb{E}[\hat{p}] = \frac{4x}{5} + \frac{2x^2}{15} + \frac{4x^3}{105} + \frac{101x^4}{21} - \frac{1252x^5}{105} + 10x^6 - \frac{20x^7}{7}.$$

However, it is possible to write a function that computes  $\mathbb{E}[\hat{p}]$  from first principles:

```
# Return expected value of the "series win percentage" estimator
GetEstimatorMean <- function(p, n_wins1, n_wins2 = n_wins1) {
  # first line is for when player 1 wins, second line is for when
  player 2 wins
  summands1 <- sapply(0:(n_wins2 - 1),
                      function(x) n_wins1 / (n_wins1 + x) * p^n_wins1 *
(1 - p)^x *
                      choose(x + n_wins1 - 1, x)
```

```

)
summands2 <- sapply(0:(n_wins1 - 1),
                    function(x) x / (x + n_wins2) * p^x * (1 -
p)^n_wins2 *
                    choose(x + n_wins2 - 1, x)
)
return(sum(c(summands1, summands2)))
}


```

Let's plot these exact values on the previous plot to make sure we got it right:

```

lines(p, sapply(p, function(x) GetEstimatorMean(x, n_wins1)),
      col = "red", lty = 2, lwd = 1.5)

```



### What might be a better estimator?

Let's go back to the setup where A and B play a best-of-7 series (or equivalently, first to 4 wins), and A wins 4 of them. Instead of estimating A's win probability by  $4/7$ , we could use the plot above to find which value of  $p$  would have given  $\mathbb{E}[\hat{p}] = 4/7$ , and use that instead. This is essentially a **method of moments estimator** with a single data point, and so has some nice properties (well, at least it's **consistent**).


The picture below depicts this method for obtaining estimates for when A wins 4 out of 4, 5, 6 and 7 games.

```

p <- 20:40 / 40
original_estimate <- c(4 / 7:4)
mm_estimate <- c(0.56, 0.64, 0.77, 1) # by trial and error
plot(p, sapply(p, function(x) GetEstimatorMean(x, n_wins1 = 4)), type =
"l",
      asp = 1, ylab = "Win prob estimator mean")
abline(0, 1, lty = 2)

```

```
segments(x0 = 0, x1 = mm_estimate, y0 = original_estimate, col = "red",
lty = 2)
segments(y0 = 0, y1 = original_estimate, x0 = mm_estimate, col = "red",
lty = 2)
```




***What are some other possible estimators in this setting?***

### **First to $n$ wins: generalizing the above**

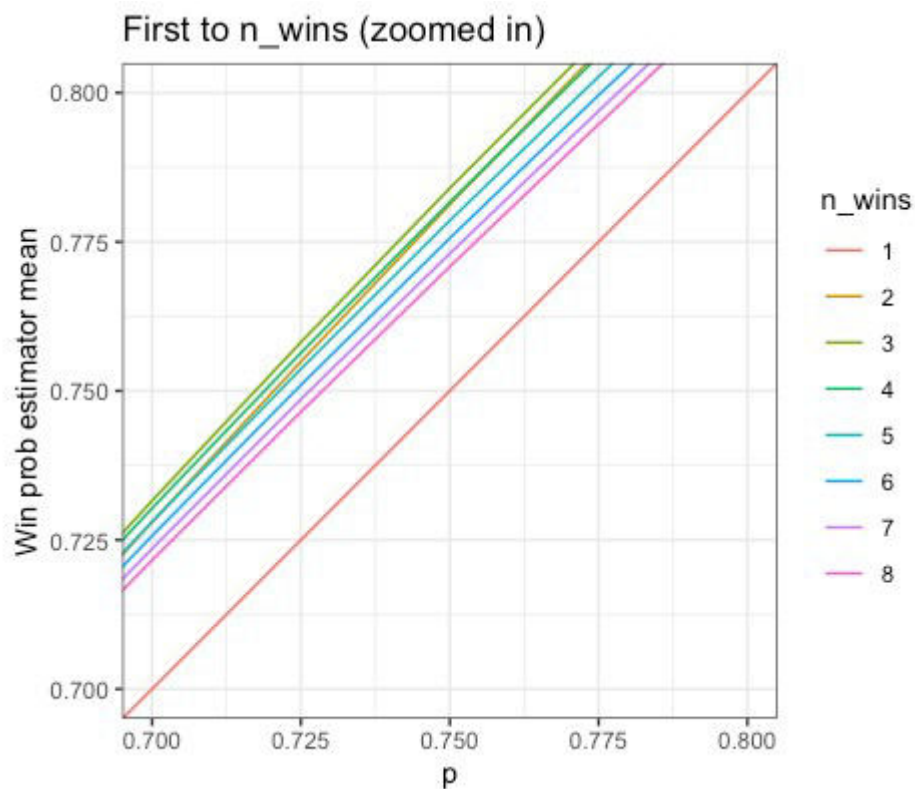
Since we've coded up all this already, why look at just first to 4 wins? Let's see what happens to the estimator  $\hat{p}$  when we vary the number of wins the players need in order to win the series.

The code below makes a plot showing how  $\mathbb{E}[\hat{p}]$  varies as we change the number of wins needed from 1 to 8 (inclusive).

```
library(tidyverse)
theme_set(theme_bw())
df <- expand.grid(p = 20:40 / 40, n_wins1 = 1:8)
df$mean_phat <- apply(df, 1, function(x) GetEstimatorMean(x[1], x[2]))
ggplot(df, aes(group = factor(n_wins1))) +
  geom_line(aes(x = p, y = mean_phat, col = factor(n_wins1))) +
  coord_fixed(ratio = 1) +
  scale_color_discrete(name = "n_wins") +
  labs(y = "Win prob estimator mean", title = "First to n_wins")
```



Past some point, the bias decreases as  $n_{\text{wins}1}$  increases. Peak bias seems to be around  $n_{\text{wins}1} = 3$ . The picture below zooms in on a small portion of the picture above:



### Asymmetric series: different win criteria for the players

Notice that in the code above, we maintained separate parameters for player 1 and player 2 ( $n_{\text{wins}1}$  and  $n_{\text{wins}2}$ ). This allows us to generalize even further, to explore what happens when players 1 and 2 need a different number of wins in order to win the series.

The code below shows how  $E[\hat{p}]$  varies with  $p$  for different combinations of  $n\_wins1$  and  $n\_wins2$ . Each line within a facet corresponds to one value of  $n\_wins1$ , while each facet of the plot corresponds to one value of  $n\_wins2$ . The size of  $\hat{p}$ 's bias increases as the difference between  $n\_wins1$  and  $n\_wins2$  increases.

```
df <- expand.grid(p = 20:40 / 40, n_wins1 = 1:8, n_wins2 = 1:8)
df$mean_phat <- apply(df, 1, function(x) GetEstimatorMean(x[1], x[2],
x[3]))
ggplot(df, aes(group = factor(n_wins1))) +
  geom_line(aes(x = p, y = mean_phat, col = factor(n_wins1))) +
  geom_abline(slope = 1, intercept = 0, linetype = "dashed") +
  coord_fixed(ratio = 1) +
  scale_color_discrete(name = "n_wins1") +
  facet_wrap(~ n_wins2, ncol = 4, labeller = label_both) +
  labs(y = "Win prob estimator mean", title = "Player 1 to n_wins1 or
Player 2 to n_wins2")
```

